



Self-observation: head-camera positions calibrate both arm chains

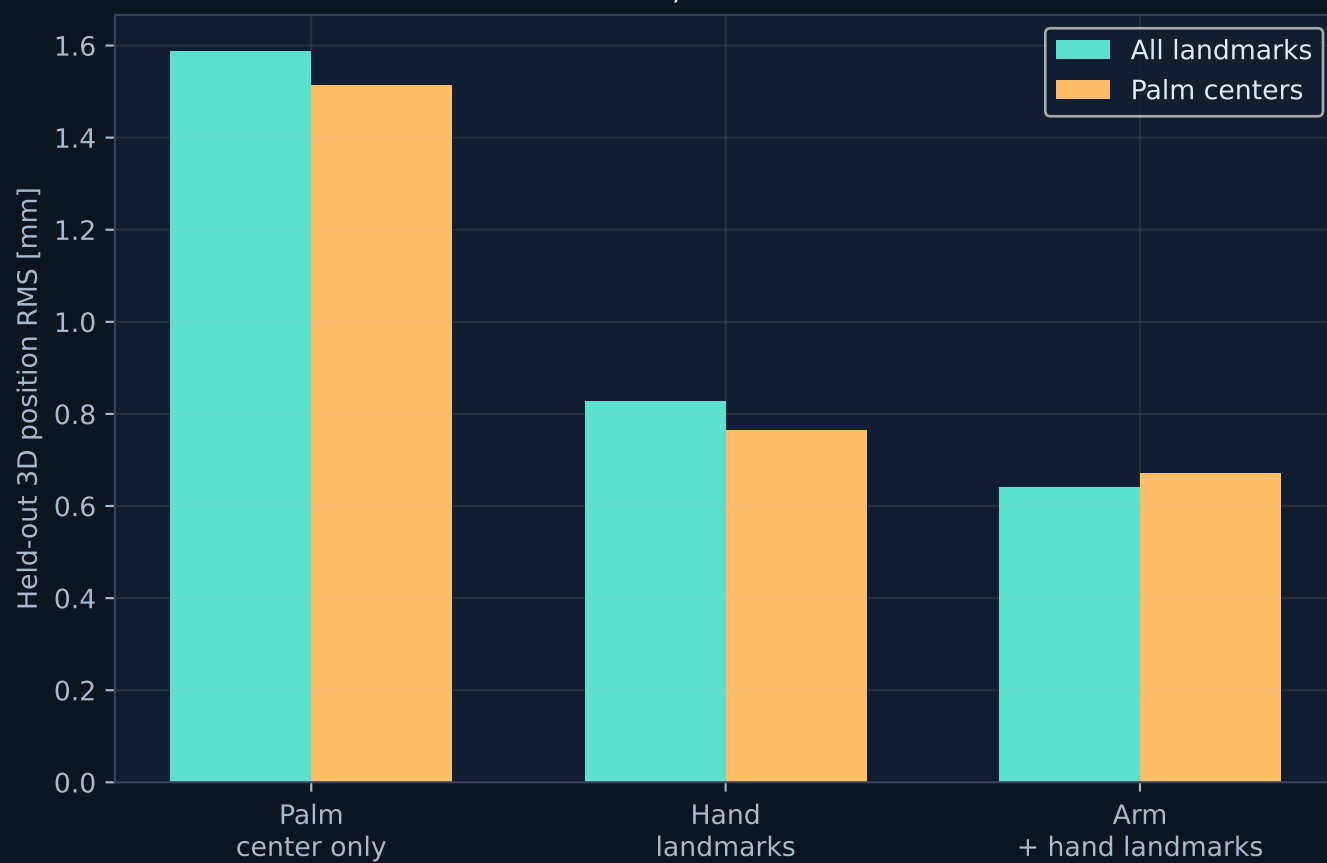
Head-camera projection of 3D landmarks



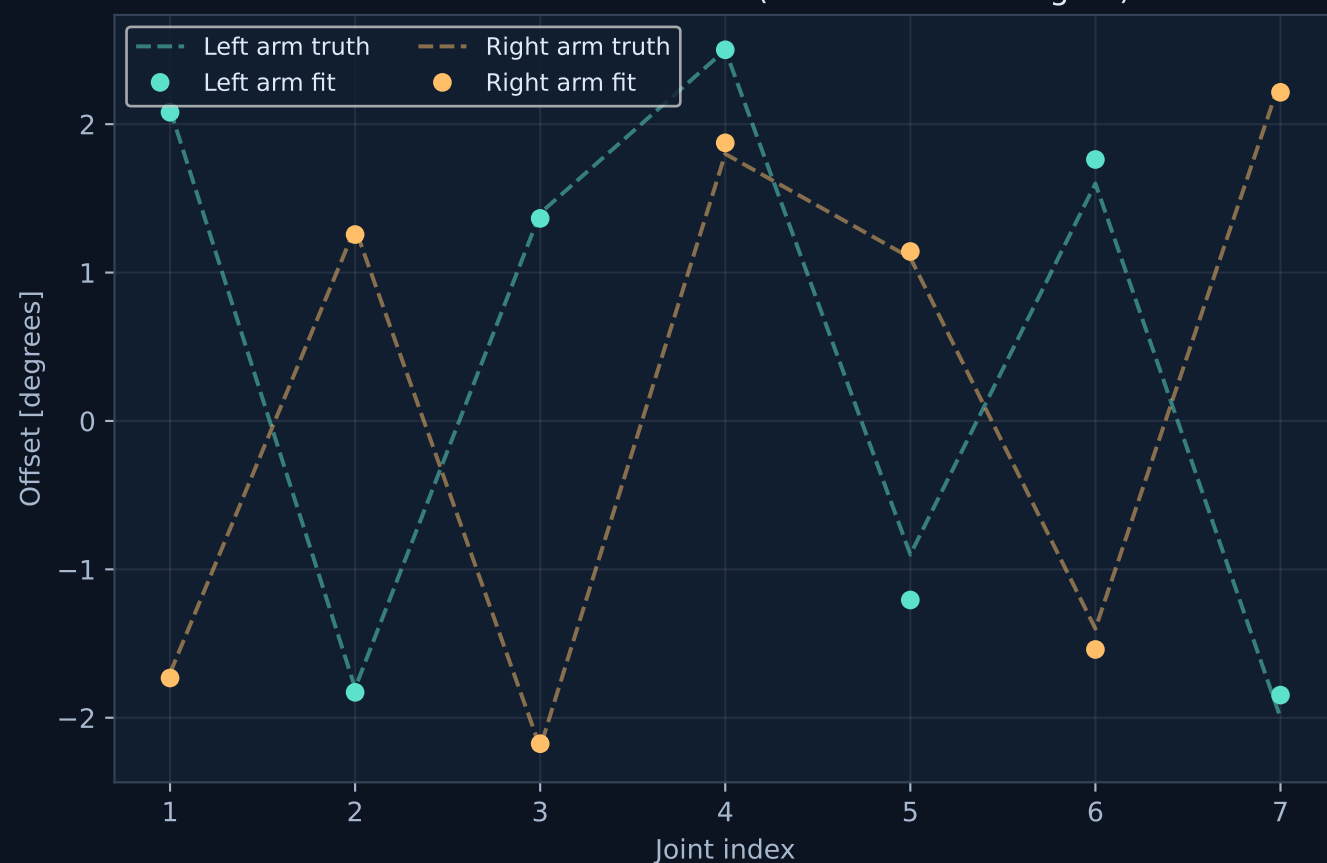
Generalization to 160 unseen configurations



Same camera frames, different landmark subsets



Recovered encoder offsets (also fits 4 link lengths)



Synthetic perception • 2 mm lateral / 4 mm depth noise • 15% dropout + 4% outliers • known camera transform • no orientation input